

Missing Layer Inference (MLI)

A Diagnostic Protocol for Structural Blind Spots in Learning-Based Decision Systems

Author: Chen-En Hou

Revised draft — 2026

Introduction

MLI is a diagnostic protocol for a specific situation: a model's own accuracy metrics look fine, but the system it feeds is getting worse. Standard debugging finds nothing because there is nothing wrong with the model in the way debugging checks for. The problem is upstream of the model — a factor that was never in the data and has no place in the architecture.

The protocol runs in three phases. Phase I rules out the two more common explanations for this pattern — bad data and a genuinely inadequate model — before assuming anything is structurally missing. Phase II turns a surviving anomaly into one or more specific, evidence-scored hypotheses about what the missing factor is. Phase III decides whether any of those hypotheses justify a real intervention, and assigns ownership, risk limits, and a rollback plan to whatever gets approved.

It is built for teams running production decision systems — supply chain, healthcare operations, financial risk, public-sector evaluation — where "the model looks fine" has stopped being a satisfying answer to "why did this get worse." It sits beside the model as an observation layer; it does not retrain it, replace it, or require access to its internals.

Origin

The idea started from a picture, not a problem statement. A standard diagram of a neural network — inputs on the left, hidden layers in the middle, outputs on the right — with a circle drawn off to one side, unlabeled, representing whatever influences the outcome but was never fed into the network.

That circle is easy to draw and easy to ignore. Most model diagnostics are built to check what's inside the box — the weights, the loss curve, the residuals against known inputs. Almost none of them are built to ask what's in that circle, because by definition nothing in the model's own telemetry can point to it. MLI is the attempt to build a process around that blind spot: not to fill the circle with certainty, but to find out, with evidence and at a defensible cost, whether it's worth looking into at all.

I. Problem Definition

Neural networks — DNN, CNN, RNN, Transformer — learn only from what is present in their training data and expressible within their architecture. A factor absent from both is invisible to the model regardless of data volume or training quality. These unmodeled factors constitute the Missing Layer. They can affect decision efficiency, system stability, and risk accumulation well before they show up as a measurable prediction error.

When MLI applies

- The model performs acceptably by its own metrics, but the system shows unexplained degradation or anomalies.
- The decision is consequential enough that internal model confidence is not a sufficient basis for trust.
- There is reason to suspect an external factor that current data or model structure has never captured.

Before MLI is invoked, two more common explanations have to be excluded first: insufficient model capability (weak architecture, wrong loss function, underfitting) and incorrect problem definition (misaligned objectives, metrics that don't track real outcomes, assumptions that no longer match operating conditions). MLI addresses what

remains once these are ruled out — it is not a first response to any anomaly.

II. Core Questions

- Q1 — Does the system contain a critical influencing factor not captured by its structure or data?
- Q2 — If so, what is it likely to be, and how much is it affecting output efficiency and stability?
- Q3 — Before that factor can be formally learned and validated, should decisions and resource allocation be adjusted now?

III. Diagnostic Triggers

1. Structural Omission

Training and test performance stay within normal range. Real-world outcomes deteriorate anyway. Standard debugging and model diagnostics find nothing.

2. Degradation Drift

Operational efficiency declines, output volatility increases, and correction costs keep rising — in a pattern the existing model offers no account of.

3. Decision Blind Spot

Frontline staff notice something is off. Management dashboards show nothing unusual. No major incident has occurred yet.

IV. Design Principle — A Method-Agnostic Three-Step Protocol

MLI is defined by a sequence of judgments, not by any specific statistical technique. The sequence is: exclude the trivial explanations, test for structure using whatever method fits the domain, then decide whether the finding justifies action. This is the same underlying logic used in differential diagnosis, root-cause investigation after an industrial accident, and security-incident triage — three fields that arrived at this pattern independently because it works. MLI applies the same discipline to AI-driven decision systems, where it has rarely been made this explicit or this systematic.

Step A — Exclude the trivial errors

Rule out bad data and an inadequate model before assuming anything is missing. This is where most apparent anomalies actually resolve. It corresponds to Phase I below.

Step B — Test for structure, with whatever tool fits

Confirm the anomaly is repeatable and structured, not noise. Phase II demonstrates this with residual clustering, counterfactual removal, and slice analysis — tools suited to a data-rich environment with an analytics team. Where those conditions don't hold, a field audit, structured interviews, or a simple A/B test can serve the same function. What carries across every version is the requirement that the method be appropriate to the evidence on hand and that repeating it would produce the same conclusion.

Step C — Decide whether it is worth treating

Finding a structural gap does not obligate the organization to fix it. Step 5.6's Value-of-Information gate is one way to make that resourcing call explicit; any comparable cost-benefit judgment serves the same function.

An organization that substitutes lighter methods in Step B, while keeping the discipline of excluding trivial causes first and gating on cost before acting, is still running MLI — the statistical toolkit in Phase II is a reference

implementation, not a requirement.

V. The MLI Pathway — Eleven Steps

Phase	Step	Purpose
I — Pre-Screening	1. Adequacy Screening	Rule out data bias or real-world misalignment before blaming the model
	2. Model Alignment	Rule out insufficient model capability or design flaws
	3. Problem Alignment	Rule out a mis-specified decision objective or problem framing
II — Diagnosis	4. Existence Assessment	Determine whether MLI activation is warranted
	5. Scope Delimitation	Build the candidate Missing Layer hypothesis space
	6. Evidence Evaluation	Score the credibility of each hypothesis
	7. Prioritization	Rank hypotheses by impact and credibility
III — Governance	8. Intervention Decision	Decide whether the finding warrants action
	9. Intervention Selection	Choose a technical, process, or governance response
	10. Risk & Accountability	Assign resources, ownership, and rollback conditions
	11. Outcome Feedback	Feed results back into evidence and future thresholds

VI. Phase I — Pre-MLI Screening

All three checks below need to pass — or be explicitly flagged and accepted as a documented organizational exception — before MLI can be activated. This phase stops teams from mistaking a world-level or model-level problem for a structural one.

Step 1. Adequacy Screening

Checks whether the data used to train and operate the model actually represent the real-world phenomena they stand in for.

1.1 Proxy Validity

Are critical real-world phenomena being represented by proxies that are distorted, convenient rather than important, or degrading over time?

- Are easily obtainable variables substituting for variables that actually matter?
- Is what's measurable diverging from what's important?
- Is the relationship between proxy and reality degrading over time?
- Are proxies being manipulated or distorted by institutional incentives or behavioral responses?

Signals: model performance stable while decision outcomes deteriorate; different data sources conflicting on the same phenomenon; a widening gap between user perception and model output; metrics improving while the underlying problem persists.

1.2 Observational Coverage

Are critical real-world conditions under-captured, or are decisions running on a stale picture of reality because of reporting delay or data drift?

- Do data voids exist within the observable decision space?
- Are low-frequency, high-impact events entirely absent from the record?
- Are specific populations, scenarios, or operating conditions systematically underrepresented?
- Do reporting delays or cross-system synchronization failures exist?

Signals: persistent failure in specific scenarios; disproportionately large errors at the edges; correction costs concentrated in under-observed regions; system responses lagging frontline conditions.

1.3 Institutional Unobservability

Do factors exist that are inherently hidden by institutional arrangements, incentive structures, or power dynamics?

- Are there behaviors that deliberately evade monitoring or audit?
- Do informal systems or shadow processes exist alongside the formal one?
- Does organizational culture obstruct information flow?
- Do incentive mechanisms distort what gets recorded in the first place?

Signals: inspection results and post-incident investigations disagreeing significantly; compliance data persistently out of step with real conditions; incidents originating from areas the data marks as safest.

1.4 Observational Adequacy — Evaluation Criteria

Category	Sub-item	Question	Method	Outcome
Proxy Validity	Proxy Drift	Does proxy–target correlation change materially over time?	Δ corr, stability index	Pass / Flag
	Temporal Drift	Is the proxy–reality relationship degrading?	Δ corr over time	Pass / Flag
	Proxy Tension	Do model metrics diverge from actual outcomes?	KPI vs. outcome gap	Pass / Flag
	Proxy Redundancy	Do better unused proxies exist?	Feature attribution	Pass / Flag
	Adversarial Shift	Are indicators being manipulated institutionally?	Expert review / field check	Pass / Flag
	Proxy Validity	Do current proxies represent the real factor?	Expert review	Pass / Flag
Observational Coverage	Coverage Gap	Do voids exist in the feature space?	Density analysis / OOD	Pass / Flag
	Rare–Impact	Are rare, high-impact events ignored?	Risk $\times$ frequency	Pass / Flag
	Error Localization	Are errors concentrated in specific segments?	Error heatmap	Pass / Flag
	Observability	Is the unobserved region acceptable in size?	Governance assessment	Pass / Flag
Institutional Unobservability	Incentive Distortion	Do incentives encourage falsification?	Institutional review	Pass / Flag
	Shadow Process	Do gray or shadow mechanisms exist?	Field investigation	Pass / Flag
	Power Blocking	Does organizational power obstruct information flow?	Organizational diagnostics	Pass / Flag

Step 2. Model Adequacy

Checks whether the model's architecture, training coverage, and stability are sufficient before treating a gap as structural rather than a modeling shortfall.

2.1 Model Capability Adequacy

- Does the model show evident underfitting or systematic error?
- Is it unable to capture known, reasonable relational structures?
- Has it received sufficient training in the critical decision scenarios?
- Does increasing capacity or changing architecture meaningfully improve performance?

Signals: persistent failure under controlled testing; substantial improvement after a capacity increase; structurally biased error distribution; elevated error rates in specific scenarios.

2.2 Objective and Temporal Alignment

- Is the loss function misaligned with the real decision risks or cost structure?
- Do evaluation metrics obscure critical risks or long-term consequences?
- Is the model aging as the environment changes?
- Does the update frequency lag behind real-world change?

Signals: metrics improving without corresponding outcome improvement; unstable decision quality in high-risk scenarios; accelerating failure in novel environments.

2.3 Decision Stability and Adversarial Resilience

- Do minor input variations cause disproportionate output changes?
- Are certain populations or scenarios highly unstable?
- Do adversarial strategies exist that can influence outputs?
- Is the model exploitable through rules, metrics, or institutional design?

Signals: similar cases yielding divergent predictions; inconsistent performance across scenarios; high sensitivity to adversarial manipulation.

2.4 Model Adequacy — Evaluation Criteria

Category	Sub-item	Question	Method	Outcome
Model Capability	Architectural Expressiveness	Can the model fit known relationships?	Error analysis / fit testing	Pass / Flag
	Training Coverage	Are decision scenarios adequately covered?	Coverage / OOD analysis	Pass / Flag
Objective & Temporal Alignment	Model Aging	Does performance degrade over time?	Performance drift analysis	Pass / Flag
	Objective Alignment	Is the loss function aligned with the decision objective?	Expert review	Pass / Flag
Decision Stability	Decision Stability	Is behavior stable under perturbation?	Sensitivity testing	Pass / Flag
	Adversarial Risk	Is the model vulnerable to manipulation?	Adversarial testing	Pass / Flag
Decision Impact	High-Risk Stability	Is behavior stable in critical scenarios?	Stress testing	Pass / Flag
	Decision Sensitivity	Do minor errors cause major consequences?	Risk assessment	Pass / Flag
Controllability	Intervenability	Can operators intervene or adjust the system?	Process review	Pass / Flag
	Failure Isolation	Can failures be contained?	Architecture review	Pass / Flag
Accountability Baseline	Traceability	Are decision sources traceable?	Audit	Pass / Flag
	Auditability	Are decisions subject to post-hoc review?	Audit	Pass / Flag

Step 3. Problem Adequacy

Checks whether the modeled problem still matches the actual decision task — whether boundaries, roles, and feasible actions are still correctly defined.

3.1 Decision Framing Alignment

- Is the system solving a prediction problem or a decision problem, and are the two being conflated?
- Are decision subject, authority, and frequency clearly specified?
- Are decision boundaries correctly defined — what falls outside the system's responsibility?
- Are multiple objectives being oversimplified into one (efficiency at the expense of stability and risk)?
- Do the available policy instruments correspond to the form of the model's outputs?

Signals: predictions that look reasonable but offer no actionable guidance, or worse, prompt inappropriate action; stakeholders disagreeing about what the system is actually solving; the same output triggering offsetting policy actions across departments.

3.2 Outcome and Label Validity

- Are labels chosen because they're available rather than important?
- Do indicators show Goodhart's Law effects — losing meaning once they become targets?
- Do indicators ignore tail risk or long-term consequences?
- Are labels distorted by administrative procedure — under-reporting, delayed reporting, selective reporting?

- Should the outcome be represented by multiple indicators rather than one crude proxy?

Signals: KPIs improving while real harm doesn't decline; the organization visibly "managing to the metric" while frontline conditions stay the same; indicators giving no early warning under high-risk conditions.

3.3 Intervention and Feedback Loop Adequacy

- Can system outputs be mapped to executable interventions?
- Does a real intervention–observation–correction loop exist, or does the system only produce reports?
- Do decision-induced behavioral changes feed back into the data-generating process, unmodeled?
- Is outcome collection sufficient to tell whether interventions actually work?
- Are decision authority and accountability aligned?

Signals: the system can be observed but not adjusted after deployment; the same problems recur and are handled ad hoc each time; responsibility and authority persistently mismatched.

3.4 Uncertainty Management Adequacy

- Are uncertainty ranges and confidence intervals explicitly communicated?
- Do decision processes allow deferred, staged, or low-consequence exploratory actions?
- Is there institutional pressure to force uncertainty into a definitive directive?
- Is "not knowing" treated as an acceptable state rather than a failure?

Signals: briefings presenting only single-point conclusions; decision-makers discovering after the fact that real uncertainty existed but was never surfaced; high-risk decisions finalized without any buffer.

3.5 Responsibility and Authority Alignment

- Does authority exist without responsibility, or responsibility without authority?
- Are risk consequences clearly attributable to a specific decision unit?
- Are critical decisions institutionally deflected to an ambiguous collective?
- Do responsibility-shifting mechanisms exist?

Signals: responsibility that cannot be localized after an incident; consensus language used to diffuse accountability on high-risk calls; a culture that avoids named ownership of decisions.

3.6 Problem Adequacy — Evaluation Criteria

Category	Sub-item	Question	Method	Outcome
Decision Structure	Problem Definition	Does the problem match the actual decision task?	Expert review	Pass / Flag
	Decision Boundary	Are decision boundaries clearly defined?	Document review	Pass / Flag
	Multi-objective Trade-offs	Are efficiency, risk, and stability explicitly addressed?	Decision audit	Pass / Flag
Outcome & Label	Label Validity	Do labels represent truly important outcomes?	KPI cross-check	Pass / Flag
	Goodhart Risk	Are indicators susceptible to manipulation?	Risk assessment	Pass / Flag
	Tail Risk Coverage	Are low-probability, high-impact events covered?	Scenario testing	Pass / Flag
	Administrative Bias	Are labels distorted by procedure?	Process audit	Pass / Flag
Intervention & Feedback	Multi-metric Completeness	Should multiple indicators replace a single proxy?	Expert review	Pass / Flag
	Actionability	Can outputs translate into executable actions?	Practical validation	Pass / Flag
	Closed-loop Capability	Does an intervention–observation–correction loop exist?	Process review	Pass / Flag
	Policy Feedback	Are policy feedback effects incorporated?	Expert review	Pass / Flag
Uncertainty & Risk Preference	Uncertainty Representation	Are confidence intervals presented?	Document review	Pass / Flag

	Tiered Decision-Making	Are staged or exploratory interventions allowed?	Process review	Pass / Flag
	Risk Preference Alignment	Are decision-makers' risk preferences aligned?	Organizational interviews	Pass / Flag
Temporal & Accountability Boundary	Temporal Alignment	Are decision, model, and policy time horizons aligned?	Process analysis	Pass / Flag
	Authority–Responsibility Alignment	Are authority and responsibility consistent?	Organizational audit	Pass / Flag
	Risk Offloading	Do responsibility-shifting mechanisms exist?	Organizational diagnostics	Pass / Flag

VII. Phase II — MLI Discovery Engine

Once Phase I clears, this phase turns a vague sense that something is wrong into one or more testable, falsifiable hypotheses about what the missing factor is.

Step 4. MLI Existence Determination

Activation requires all Phase I items to pass, except items explicitly flagged and accepted as an organizational exception, and at least two of the following five conditions:

Category	Indicator	Assessment Basis
Performance Gap	$\Delta\text{Perf} > \text{threshold}$	Field performance vs. model prediction
Structural Residual	Non-random residuals	Residual clustering / patterned distribution
Decision Friction	Rising decision cost	Delays, conflicts, remedial actions
Behavioral Drift	Abnormal usage behavior	Operator behavioral shifts
Correction Failure	Repeated adjustments ineffective	Model tuning / institutional changes fail

Step 5. MLI Scope Identification

Existence of a Missing Layer is inferred only when an anomaly is repeatable, structured, material, and not explained by any known cause — all four conditions together, not any one alone.

5.1–5.2 Anomaly Definition

The anomaly set A requires at least one input from each of: field performance/efficiency/stability indicators, decision cost indicators, and model or decision outputs. Example entries: median delivery delay up 20% under identical inputs; error rate of a specific population consistently $2\times$ the overall average; handling cost rising while model metrics stay flat. If A cannot be defined explicitly, the process stops here.

5.3 Repeatability Test

At least one of: the anomaly persists across different time windows, different samples or slices, or different regions or units. Decision rule: occurrence rate $\geq p$ (e.g., 60%) across at least T consecutive windows, or the same directional deviation across at least m of K slices.

5.4 Structured Residual Test

At least two of the following: residual clustering (errors concentrate in specific scenarios or nodes), residual patterning (a persistent directional bias), residual association (correlation with a recorded but unmodeled variable), or residual migration (correcting one area causes errors to reappear elsewhere — a classic sign of structural omission). The random-error hypothesis must be rejected before proceeding.

5.5 Exclusion of Known Causes

Rather than re-running all of Phase I, each anomalous slice gets a focused re-check against the relevant items only.

Group	Check Item	Question
A-Check	Proxy validity	Are the proxies still appropriate in this slice?

	Observational coverage	Is data coverage missing or degraded here?
	Institutional bias	Are institutional or behavioral distortions affecting this slice?
B-Check	Model structure	Does this slice show signs of underfitting?
	Training coverage	Is this slice outside the training distribution?
	Stability	Is model behavior highly unstable here?
C-Check	Problem framing	Does this slice exceed the original problem boundary?
	Metrics validity	Is this slice described by misleading metrics?
	Intervention feasibility	Does this slice lack any actionable intervention?

If at least one slice passes A-Check, B-Check, and C-Check simultaneously while the anomaly persists, Q1 is likely to hold and the process moves to counterfactual testing.

Counterfactual Stress Test

For each anomalous slice, run at least one counterfactual: a policy counterfactual (simulate an alternative strategy), a process counterfactual (bypass a specific node), or a resource counterfactual (modify allocation). An anomaly that survives multiple counterfactuals strengthens the Missing Layer hypothesis; one that disappears indicates a decision-use error, not a missing layer.

Value-of-Information Gate

$$VoI = (DIP \times PR \times U) - IC$$

VoI score	Action	
≥ 40	Enter Step 6 — high priority	
25–39	Enter Step 6 — normal priority	
10–24	Low-cost exploration only	
< 10	Terminate; return to monitoring	
Metric	Score	Definition
Decision Improvement Potential (DIP)	5 / 4 / 3 / 2 / 1	Prevents systemic failure / significantly reduces risk / improves efficiency / minor optimization / negligible
Probability of Occurrence (PR)	5 / 4 / 3 / 2 / 1	Repeatedly observed / many contexts / some contexts / occasional / rare
Exploration Cost (IC)	5 / 4 / 3 / 2 / 1	Large redesign / long-term commitment / cross-department coordination / low-cost experiment / negligible
Urgency (U)	5 / 4 / 3 / 2 / 1	Immediate risk / near-term / mid-term / long-term / no time pressure

Illustration: Three Real-World Cases

The following three cases illustrate what Phase II actually produces once the Detection Core runs to completion.

Case 1 — Night-Shift Failure in Supply Chain Operations

Over six months, a supply-chain system saw delivery efficiency decline 18% while model prediction error stayed within ±2%. Conventional debugging and process audits found nothing.

Case 2 — Emergency Department Triage Degradation

Patient boarding time rose 21% during holidays and night shifts while the patient-volume model maintained accuracy above 95%. Management dashboards looked normal; frontline staff consistently reported growing congestion.

Case 3 — Distortion in a Government Startup Subsidy Program

An evaluation model returned high performance predictions for three straight years while actual industrial impact failed to materialize and post-hoc correction costs rose annually.

Cross-Case Analysis

Component	Case 1	Case 2	Case 3
Anomaly Set A	Order delay +18%, return rate +9%	ED boarding +21%	Industrial impact unobserved

Structural Slices S	Night-shift logistics, third-party drivers	Night shifts, holidays	Startup firms
Structural Residuals R	No model version explains it	Predictions remain normal	Evaluation scores stay high
Counterfactual Tests C	Persists after removing weather & volume	Persists after removing patient flow	Persists after excluding funding factors
VoI	14.6	11.2	12.8
Missing Layer Description	Night-shift driver dispatch decisions unmodeled	Triage decisions deviate from design assumptions	Founders' execution capability unmodeled
Impact Mechanism	Manual dispatch creates bottlenecks	Misclassification of patient severity	Execution gap
Impact Scope	Delivery efficiency, complaints	Bed utilization	Program effectiveness
Falsification Condition	No improvement after inclusion	No improvement in any indicator	No significant difference

Resulting hypotheses:

- h1 — The actual dispatch decisions of third-party night-shift drivers are unrepresented in the current supply-chain model and produce structural impact on system performance.
- h2 — Night-shift medical staff's actual triage decisions deviate substantially from system design assumptions, and this deviation is uncaptured by any model or process.
- h3 — The real execution capability of founders and core teams is absent from the subsidy evaluation model and constitutes a critical missing layer driving outcome distortion.

All three enter Step 6 for evidence evaluation.

Step 6. MLI Evidence Evaluation

Step 5 establishes that a Missing Layer is likely. Step 6 answers a separate question: whether a specific hypothesis carries enough evidence and credibility to reach the decision layer. It converts a plausible explanation into an evidence base a decision-maker can actually stand behind.

Evidence Types

Evidence Type	Description	Evaluation Criteria
Structural	Logical gaps in architecture, process, or institutional design	Explicit unmodeled structure (Yes/No)
Behavioral	Abnormal behavior of users, staff, management	Cross-role consistency (High/Med/Low)
Statistical	Residuals, error distributions, degradation curves	Residual deviation $\geq 2\sigma$ = strong
Counterfactual	Whether the phenomenon disappears if the hypothesis holds	Anomaly disappears under hypothesis (Yes/No)
External	External cases, cross-system comparison, expert opinion	External case alignment (High/Med/Low)

A hypothesis that explains only one evidence source is treated as low credibility regardless of how compelling that single source looks.

Evidence Strength Scoring

Evaluation Item	0 pts	1 pt	2 pts
Evidence Volume	One source, not firsthand	≥ 2 independent sources, one indirect	≥ 2 firsthand sources, one with auditable records
Evidence Consistency	Evidence conflicts	Mostly supports, some conflict remains	Highly consistent, no material conflict
Explanatory Power	Explains isolated anomalies only	Explains several, major gaps remain	Explains the core anomaly set and its evolution
Alternative Exclusion	No systematic exclusion	Partial exclusion of major alternatives	All known major alternatives excluded, documented
Verifiability	No feasible verification	Verifiable but costly or slow	Verifiable within acceptable cost, risk, time

Spurious Correlation and Alternative Explanation Checklist

- Short-term incidental phenomenon? Each "Yes" subtracts one point from the total.

- External shock event?
- Data bias?
- Model defect? Three or more "Yes" defers the hypothesis; all four excludes it.

Total Score	Credibility	Action
8–10	High	Proceed to Step 7
5–7	Medium	Collect more evidence, reassess
0–4	Low	Terminate
Excluded	—	Archive

Alternative-exclusion credit (the 2-point tier above) requires documenting at least two of: cross-validation against known model defects, against data-bias hypotheses, against external shocks, and against known structural factors. Without documentation, that item is capped at one point. Verifiability credit similarly requires a testable design (experiment, A/B test, backtest, simulation, sandbox), completion within a decision-acceptable timeframe, cost under 30% of the potential risk impact, and no irreversible damage to the core system from the test itself — otherwise the item is capped at one point.

Step 7. MLI Prioritization

High- and medium-credibility hypotheses are ranked by systemic impact and remediation cost, so limited resources go to the highest-value targets.

$$\text{Priority Index} = (I \times P \times D \times T) \div C$$

Score	Impact Severity (I)	Probability (P)	Diffusion (D)	Treatability (T)	Remediation Cost (C)
1	Single-metric impact	<1%/period	Single node	Nearly non-intervenable	<1 person-week
2	Local performance impact	1–5%/period	Adjacent processes	High difficulty	1–4 person-weeks
3	Cross-process impact	5–20%/period	Multiple processes	Intervenable, high cost	1–3 person-months
4	Systemic / regulatory impact	20–50%/period	Cross-system	Feasible	3–12 person-months
5	System failure / major incident	>50%/period	Cascading failure	Easily/rapidly intervenable	>12 person-months

Priority levels: 60+ Critical · 30–59 High · 10–29 Medium · below 10 Low.

VIII. Phase III — Decision and Governance

Here a Missing Layer stops being an analytical finding and becomes something the organization must resource, own, and answer for.

Step 8. Intervention Decision Gate

Dimension	Question	Yes/No
Evidence Maturity	Is Step 6 credibility ≥ High?	☐ / ☐
Governance Priority	Is Step 7 priority ≥ High?	☐ / ☐
Risk Necessity	Is the risk unacceptable without intervention?	☐ / ☐
Solution Feasibility	Are executable options available?	☐ / ☐
Accountability Capacity	Are owners and resources clearly assigned?	☐ / ☐

All five conditions must clear before proceeding to Step 9. Any single "No" means the item is registered as a governance monitoring item and reassessed later, not acted on.

Step 9. Intervention Mode Selection

Each surviving hypothesis is translated into a concrete plan — what to modify, what to instrument, which processes to adjust — and reduced to a Minimum Viable Intervention (MVI).

Case 1 — Third-Party Night-Shift Driver Dispatch

Gap: real dispatch data (overrides, late acceptance, no-shows, route deviations) is missing; driver behavioral variability and night-shift institutional differences are unmodeled.

Type	Decision	Deliverable
------	----------	-------------

Data	Complete night-shift driver event logs (accept/refuse/cancel/timeout/reroute)	Event schema + ETL + daily completeness report
Model	Introduce driver availability/reliability features	Feature definitions + baseline experiments
Process	Require override logging in night-shift dispatch	Override fields + mandatory rules
Governance	Align third-party KPIs/contracts with night-shift risk	KPI clauses + monitoring dashboard

MVI-1: data and process first — build the event log and mandatory override recording, so the missing layer becomes observable before attempting a full fix. Completion criteria: event completeness $\geq 90\%$, override reason coverage $\geq 80\%$, anomaly set A decomposed into 2–3 reproducible patterns across slice S.

Case 2 — Night-Shift Triage Deviation

Gap: actual triage rule usage, night-shift resource constraints, and implicit trade-offs (boarding speed vs. complaint avoidance) are unmodeled.

Type	Decision	Deliverable
Data	Triage reason codes + revision trail	Log schema + revision markers
Process	Night-shift triggered secondary review	Review thresholds + SOP
Model	Add night-shift constraint features	Constraint features + benchmark
Governance	Shift KPIs from speed to correctness $\times$ safety	KPI revision + exception handling

MVI-2: build the triage audit trail and triggered secondary review before touching the model. Completion criteria: reason coverage $\geq 85\%$, revision traceability $\geq 80\%$, at least one actionable process point identified for the congestion anomaly.

Case 3 — Founder Execution Capability

Gap: execution proxies (delivery history, milestone punctuality, team stability) are missing; incentive distortion mistakes writing quality for execution ability.

Type	Decision	Deliverable
Data	Milestone punctuality / delivery quality history	Milestone table + data dictionary
Process	Staged funding + gate reviews	Gate design + disbursement rules
Model	Execution capability proxy indicators	Proxy indicators + weighting scheme
Governance	Define execution risk as a veto condition	Veto rules + evidence requirements

MVI-3: staged funding gates plus milestone datafication — make execution ability traceable before attempting to model it. Completion criteria: milestone coverage $\geq 90\%$, gate reviews institutionalized (e.g., Month 2 / Month 4 checkpoints), at least one data pattern separating high-score failures from high-score successes.

Step 10. Intervention Risk and Accountability

Selecting an MVI is not the same as being authorized to run it. This step assigns explicit probability, impact, rollback difficulty, and accountability level before anything is executed.

Score	Probability (P)	Impact (I)	Rollback Difficulty
1	<5%, no precedent	Negligible, no reporting required	Immediate rollback, no loss
2	5–15%, occasional	Local, recoverable within unit	Low cost, recovery within days
3	15–40%, seen in similar cases	Cross-department, needs management	Moderate cost, recovery in weeks
4	40–70%, frequent	Major operational/safety/reputational	High cost, rollback itself risky
5	>70%, near-certain	Regulatory/life-safety/organizational survival	Nearly irreversible

$$\text{Risk Level (R)} = P \times I$$

R	Level	Rule
1–4	Low	Executable, unit-level control
5–9	Medium	Supervisor approval required
10–15	High	Senior management / committee approval required
16–25	Extreme	Prohibited without highest-authority approval

Accountability tiers: A — Departmental (manager); B — Cross-Departmental (senior management / committee); C — Institutional (chief responsible authority); D — Regulatory (external, life-safety, or public-safety liability, requiring prior approval).

Case	Primary Risk	P	I	R	Rollback Risk	Resources
1. Night-shift drivers	Third-party misreporting distorts diagnosis	4	3	12 (High)	3	Ops 1 + Data/IT 1, 6–8 weeks
2. ED triage	Slower process / mis-triage patient-safety	3	5	15 (High)	4	ED clinical 1 + IT 1, 8–10

	incident					weeks
3. Subsidy evaluation	Administrative burden causes execution stagnation	3	3	9 (Medium)	2	PMO + reviewers, 1 subsidy cycle

Rule: intervention proceeds if $R \leq 15$ and rollback risk ≤ 3 ; otherwise it escalates to a higher accountability tier or is terminated.

Step 11. Outcome Feedback

$$\text{Intervention Value} = \Delta \text{Outcome} - \Delta \text{Risk} - \Delta \text{Cost}$$

Dimension	Metric	Threshold	Feedback Action
$\Delta \text{Outcome}$	(After – Before) / Before	$\geq +10\%$	Increase MLI evidence strength
ΔRisk	R before – R after	≥ 3	Confirm intervention effectiveness
ΔCost	Actual – Estimated	$\leq +10\%$	Within control
Incidents	Major event count	0	Maintain strategy
User Feedback	Survey / interview positivity	$\geq 70\%$	Support hypothesis
Condition		Action	
$\Delta \text{Outcome} \geq 10\%$ and $\Delta \text{Risk} \geq 3$		Institutionalize the correction	
$\Delta \text{Outcome} < 5\%$ or $\Delta \text{Risk} < 1$		Reduce MLI weight	
New risks emerge		Return to Steps 5–6	
Intervention fails		Trigger rollback, downgrade MLI	

IX. Where MLI Sits in the System

MLI operates as a post-decision, cross-layer observation point. It reads a standardized event stream — input summary, model output, decisions and actions taken, delayed outcome feedback, and residual/anomaly metrics — without modifying the underlying model. It is most useful just before a high-impact decision, when anomalies are accumulating, or when early warning signals appear before they've become a measurable incident.

Data Contract

A minimal universal event schema (JSON / Avro / Protobuf): system_id, model_id, model_version, timestamp, slice_keys, input_fingerprint (hash or embedding, not raw input), output (score/label + confidence), decision (action, policy, override flag), outcome (delayed KPI, ground truth, incidents), diagnostics (drift, latency, resource usage).

Deployment

A sidecar service collecting events and aggregating feedback is the lowest-friction option and the one most organizations adopt first. A plugin SDK (Python for ML workloads, Java/Go for backend services, JS/TS for front-end and workflow) covers direct event publishing where a sidecar isn't practical. A dedicated MLI assist model — for hypothesis generation, evidence scoring, or intervention recommendation — is an optional later addition; it consumes only the standardized event schema and stays decoupled from the underlying production model.

Compatibility Risks

- Inconsistent slice fields across systems — mitigated with a configurable slice_keys dictionary and a minimal common set.
- Delayed feedback and sparse labels — mitigated with built-in delay windows (7/30/90 days) and provisional proxy outcomes updated once true labels arrive.
- Data privacy and access control — mitigated by storing fingerprints only, hashing sensitive fields, and keeping the evidence chain auditable without exposing plaintext.

Acknowledgement

CC BY-NC 4.0 © 2026 Hou, Chen-En. Licensed under Creative Commons Attribution-NonCommercial 4.0 International. Non-commercial sharing and adaptation permitted with attribution; commercial use requires written permission.